

# Homework 3, Part II Learning, Inference, Decisions

Nikhil Pereira nmp54

April 18th, 2022

## Problem 4

### Part A - Download the Data and Load into Pandas

```
In [17]: ▶ import pandas as pd
import warnings
warnings.filterwarnings("ignore")
```

```
In [18]: ▶ red = pd.read_csv('winequality-red.csv', sep=';')
red['Color'] = 1
white = pd.read_csv('winequality-white.csv', sep=';')
white['Color'] = 0
name_mapper = {'fixed acidity':'fixed_acidity', 'citric acid':'citric_acid', 'total sulfur dioxide': 'total_sulfur_dioxide', 'volatile acidity': 'volatile_acidity', 'free sulfur dioxide': 'free_sulfur_dioxide', 'total phenols': 'total_phenols', 'polyphenols': 'polyphenols', 'sulphates': 'sulphates', 'alcohol': 'alcohol'}
df = pd.concat([red, white], axis=0, ignore_index=True).rename(columns=name_mapper)
indices = list(df.index)
```

### Part B - Random Sample for Train Test Split

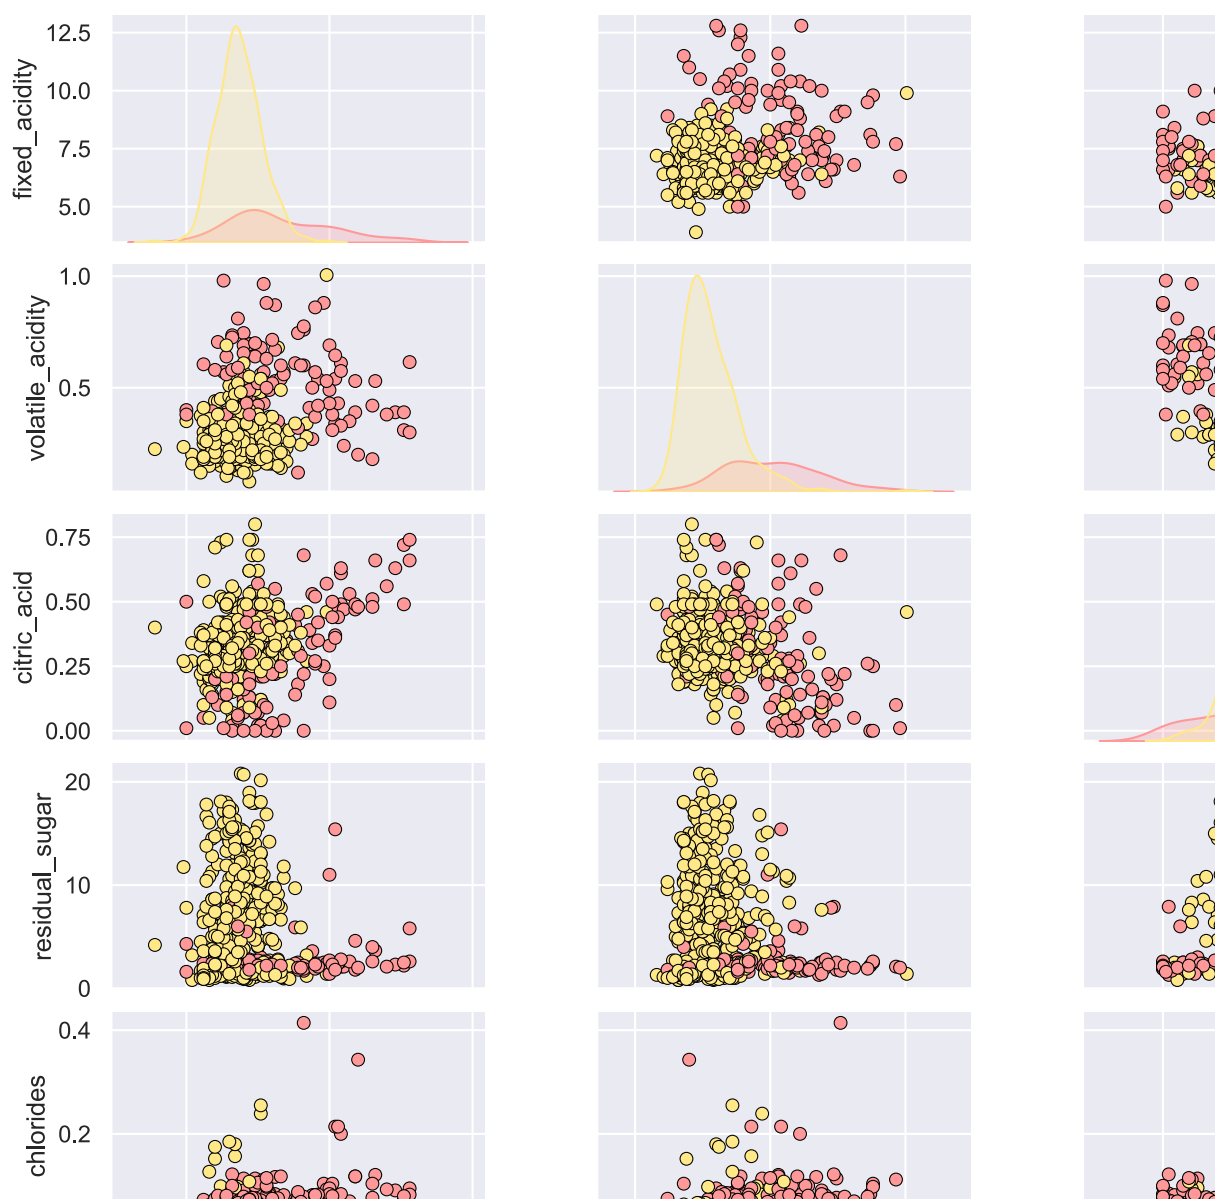
```
In [19]: ▶ from random import sample
train_indices = list(sample(list(df.index), 5000))
test_indices = list(set(list(df.index)) - set(train_indices))
train_set = df.iloc[train_indices,: ]
test_set = df.iloc[test_indices,: ]
```

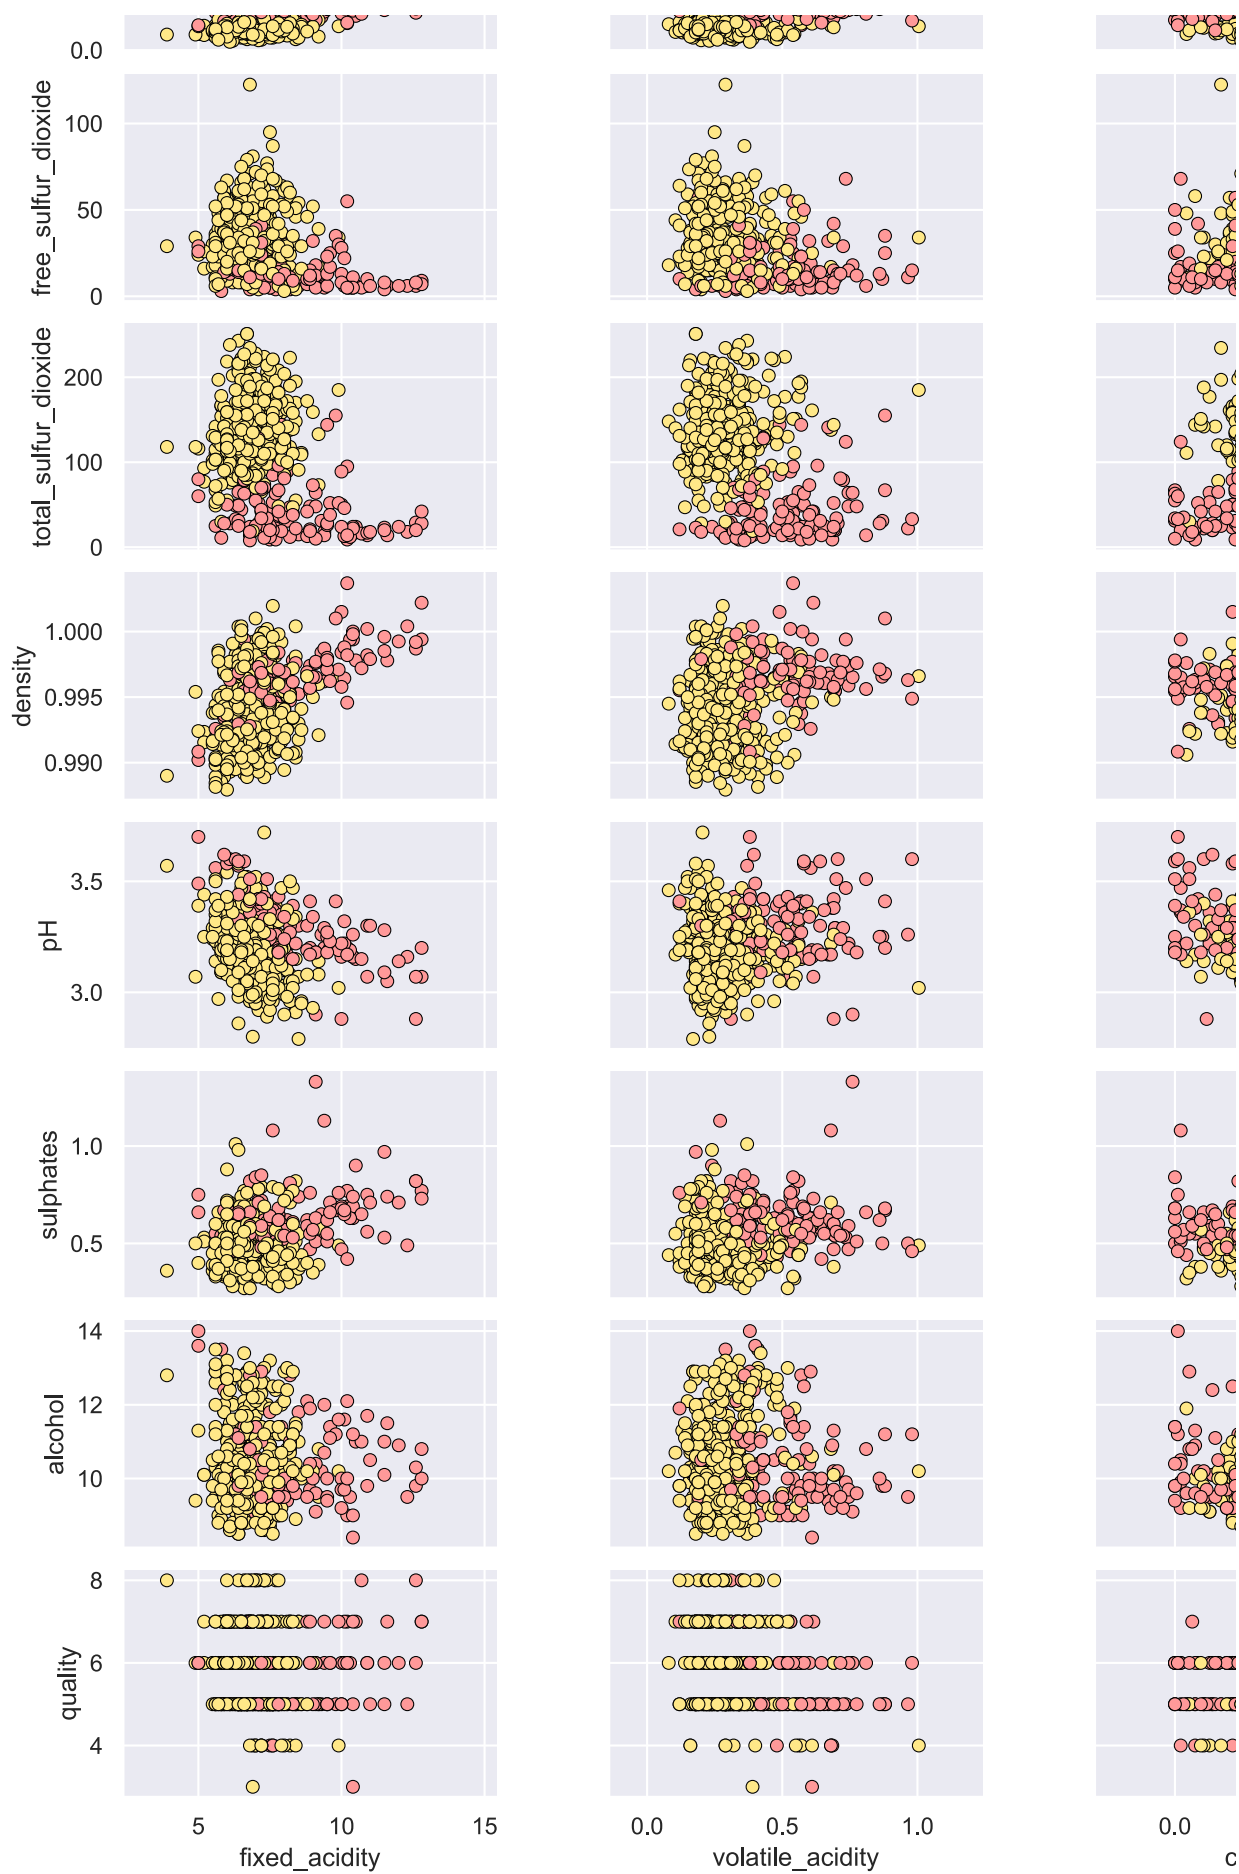
### Part C - Visualizing the Wine Set All Covariates

```
In [11]: %pylab
%matplotlib inline
%config InlineBackend.figure_format = 'svg'
import seaborn as sns
sns.set()
wt = {0:'white',1:'red'}
smaller_train_set = train_set.sample(n=500)
smaller_train_set['wine_type'] = smaller_train_set.Color.map(wt)
cols = ['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
        'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
        'pH', 'sulphates', 'alcohol', 'quality', 'wine_type']
pp = sns.pairplot(smaller_train_set[cols], hue='wine_type', height=1.8, aspect=1.8,
                 palette={"red": "#FF9999", "white": "#FFE888"},
                 plot_kws=dict(edgecolor="black", linewidth=0.5))
fig = pp.fig
fig.subplots_adjust(top=0.93, wspace=0.3)
t = fig.suptitle('Wine Attributes Pairwise Plots', fontsize=14)
```

Using matplotlib backend: Qt5Agg

Populating the interactive namespace from numpy and matplotlib





## Analyzing Important Covariates to be Transformed

```
In [5]: ▶ for c in train_set.columns:
        print('Column:', c, ' Range:', round(train_set[c].max()-train_set[c].min(),2), ' Mean: ', round(
Column: fixed_acidity Range: 11.7 Mean: 7.22 Std: 1.32
Column: volatile_acidity Range: 1.5 Mean: 0.34 Std: 0.16
Column: citric_acid Range: 1.0 Mean: 0.32 Std: 0.14
Column: residual_sugar Range: 31.0 Mean: 5.4 Std: 4.69
Column: chlorides Range: 0.46 Mean: 0.06 Std: 0.03
Column: free_sulfur_dioxide Range: 288.0 Mean: 30.58 Std: 18.06
Column: total_sulfur_dioxide Range: 434.0 Mean: 116.03 Std: 56.73
Column: density Range: 0.02 Mean: 0.99 Std: 0.0
Column: pH Range: 1.29 Mean: 3.22 Std: 0.16
Column: sulphates Range: 1.72 Mean: 0.53 Std: 0.14
Column: alcohol Range: 6.9 Mean: 10.49 Std: 1.19
Column: quality Range: 6 Mean: 5.82 Std: 0.88
Column: Color Range: 1 Mean: 0.24 Std: 0.43
```

### Comments on Scatter Matrix

Some notable positive correlations are density vs fixed acidity, density vs sulphates whereas notable negative correlated variables are fixed acidity vs pH and citric acid vs pH which makes sense (lower pH is less acidic). Free\_sulfur\_dioxide and total\_sulfur\_dioxide are two covariates that should be linearly transformed since their range and standard deviations are very high indicating a poor regression if no changes are made.

## Part D - Fitting a Linear Model on All Covariates to establish a baseline model

```
In [20]: ▶ import statsmodels.stats.api as sms
import statsmodels.api as sm
import statsmodels.formula.api as smf
linear_formula = 'quality ~' + '+'.join(list(set(df.columns)-{'quality'}))
lm = smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 Baseline: ', round(lm.rsquared,3))
```

In Sample R2 Baseline: 0.291

```
In [21]: ## Performing CV and Measuring Out Sample R^2
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.model_selection import train_test_split
import numpy as np

#Function that uses 5 fold cross validation by creating random train test splits at a 20 % thresh
# and then fits an ols model and measures the r_2 and MSE between the y_true and y_predicted in
# Finally it returns the means of r_2 and MSE
def CV_R2(features,linear_formula,data,split):
    #Create r_2 and MSE vectors
    r_2 = np.zeros(split)
    ms = np.zeros(split)
    for i in range(split):
        #Split the data into 5 folds
        X_train, X_test, y_train, y_test = train_test_split(
            data.loc[:,features], data.quality, test_size=0.2, random_state=i)
        full_train = pd.concat([X_train,y_train],axis=1)
        #Fit Model on train set
        model = smf.ols(formula=linear_formula, data=full_train).fit()
        #Gather model predictions from test set
        predictions = model.predict(X_test)
        #Calculate the R_2 and MSE
        r_2[i] = r2_score(y_test,predictions)
        ms[i] = mean_squared_error(y_test, predictions)
    #Return the mean r_2 and mean MSE found
    return r_2.mean(), ms.mean()

features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation Baseline ', CV_R2(features,linear_formula,train_set))
print('Out of Sample MSE via Cross Validation Baseline', CV_R2(features,linear_formula,train_set,

Out of Sample R^2 via Cross Validation Baseline 0.28760073577986595
Out of Sample MSE via Cross Validation Baseline 0.5356495740401874
```

## Part E - Linear Transformation for features from Baseline Model

```
In [22]: transformed = train_set.copy()
#Log transform the sulfur dioxide columns as they have very large ranges
transformed['free_sulfur_dioxide'] = np.log(transformed['free_sulfur_dioxide'])
transformed['total_sulfur_dioxide'] = np.log(transformed['total_sulfur_dioxide'])

linear_formula = 'quality ~' + '+'.join(list(set(df.columns)-{'quality'}))
lm2 = smf.ols(linear_formula, data=transformed).fit()
print('In Sample R2 after linear transform: ', round(lm2.rsquared,3))

In Sample R2 after linear transform: 0.302
```

```
In [23]: features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation on Transformed Sulfur Dioxide Columns ', CV_R2(features,linear_formula,train_set))
print('Out of Sample MSE via Cross Validation on Transformed Sulfur Dioxide Columns', CV_R2(features,linear_formula,train_set,

Out of Sample R^2 via Cross Validation on Transformed Sulfur Dioxide Columns 0.29747832578390015
Out of Sample MSE via Cross Validation on Transformed Sulfur Dioxide Columns 0.5282591535172012
```

After transforming the sulfur dioxide columns and performing the same procedure as in part D, the r-squared went up a percentage point and the mean squared error became smaller by very little. This is not a significant improvement.

## Part F - Fitting a linear model with 2-Way Interaction Terms (Report In Sample and Out Sample R^2)

```
In [27]: import itertools
#add two way interactions to the linear formula
linear_formula = 'quality ~' + '+' .join(list(set(df.columns)-{'quality'}))+['%s:%s'%v for v in ite
lm3 = smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 with 2-Way Interaction Terms: ', round(lm3.rsquared,3))
```

In Sample R2 with 2-Way Interaction Terms: 0.357

```
In [28]: features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation with 2_way Interaction ', CV_R2(features,linear_for
print('Out of Sample MSE via Cross Validation with 2_way Interaction ', CV_R2(features,linear_for
```

Out of Sample R<sup>2</sup> via Cross Validation with 2\_way Interaction -0.04774341757190788  
 Out of Sample MSE via Cross Validation with 2\_way Interaction 0.7828630408021304

After adding interaction terms the in sample R<sup>2</sup> improved by 6 percentage points but the out of sample r-squared and MSE declined as the model has more complexity and started to overfit.

## Part G - Stepwise regression with AIC criterion based on 2-Way Interaction Model Report Estimates of Prediction Error

```
In [29]: singlefeatures = list(set(df.columns)-{'quality'})
interactions = ['%s:%s'%v for v in itertools.combinations(singlefeatures,2)]
```

### Forward Stepwise Regression

```
In [30]: def fitmodel(S,data):
    return smf.ols('quality ~ '+'+'.join(S) if len(S)>0 else '1'), data=data).fit()
## forward stepwise
Sfwd = set()
# features = set(singlefeatures)
features = set(singlefeatures).union(set(interactions))
while len(Sfwd)<len(features):
    f = max(features - Sfwd, key = lambda f: fitmodel(Sfwd.union({f}),train_set).aic)
    after = fitmodel(Sfwd.union({f}),train_set).aic
    before = fitmodel(Sfwd,train_set).aic
    if after > before:
        Sfwd = Sfwd.union({f})
    else:
        break
print('Forward Stepwise Terms:', Sfwd)
```

Forward Stepwise Terms: {'residual\_sugar:alcohol', 'fixed\_acidity:sulphates', 'density:pH'}

### Backwards Stepwise Regression

```
In [32]: > ## backward stepwise
# features = set(singlefeatures)
features = set(singlefeatures).union(set(interactions))
Sbwd = set(features)
while len(Sbwd)>0:
    f = max(Sbwd, key = lambda f: fitmodel(Sbwd-{f},train_set).aic)
    after = fitmodel(Sbwd-{f},train_set).aic
    before = fitmodel(Sbwd,train_set).aic
    if after - before > 5:
        Sbwd = Sbwd-{f}
    else:
        break
print('Backwards Stepwise Terms:', Sbwd)
```

Backwards Stepwise Terms: {'fixed\_acidity:residual\_sugar', 'total\_sulfur\_dioxide:pH', 'alcohol', 'residual\_sugar:volatile\_acidity', 'pH', 'volatile\_acidity:chlorides', 'citric\_acid:Color', 'residual\_sugar:pH', 'fixed\_acidity', 'residual\_sugar:Color', 'Color', 'total\_sulfur\_dioxide:residual\_sugar', 'chlorides', 'fixed\_acidity:sulphates', 'density:Color', 'fixed\_acidity:free\_sulfur\_dioxide', 'free\_sulfur\_dioxide:alcohol', 'free\_sulfur\_dioxide:volatile\_acidity', 'free\_sulfur\_dioxide', 'density:alcohol', 'density:volatile\_acidity', 'citric\_acid:volatile\_acidity', 'total\_sulfur\_dioxide:chlorides', 'sulphates', 'residual\_sugar:alcohol', 'sulphates:alcohol', 'sulphates:chlorides', 'sulphates:volatile\_acidity', 'fixed\_acidity:alcohol', 'fixed\_acidity:total\_sulfur\_dioxide', 'alcohol:Color', 'citric\_acid:total\_sulfur\_dioxide', 'sulphates:pH', 'citric\_acid:residual\_sugar', 'total\_sulfur\_dioxide:density', 'fixed\_acidity:density', 'citric\_acid', 'fixed\_acidity:volatile\_acidity', 'total\_sulfur\_dioxide', 'density:pH', 'citric\_acid:free\_sulfur\_dioxide', 'citric\_acid:density', 'citric\_acid:sulphates', 'sulphates:residual\_sugar', 'residual\_sugar', 'free\_sulfur\_dioxide:chlorides', 'volatile\_acidity', 'citric\_acid:chlorides', 'alcohol:pH', 'Color:chlorides', 'total\_sulfur\_dioxide:Color', 'density:chlorides', 'total\_sulfur\_dioxide:alcohol', 'residual\_sugar:chlorides', 'free\_sulfur\_dioxide:density', 'total\_sulfur\_dioxide:volatile\_acidity', 'fixed\_acidity:chlorides', 'free\_sulfur\_dioxide:Color', 'free\_sulfur\_dioxide:sulphates', 'free\_sulfur\_dioxide:pH', 'sulphates:Color', 'free\_sulfur\_dioxide:residual\_sugar', 'density:sulphates', 'alcohol:chlorides', 'pH:volatile\_acidity'}

### Considering the optimal AIC from Forward and Backward Stepwise

```
In [33]: > ## stepwise
Sboth = Sfwd if fitmodel(Sfwd,train_set).aic > fitmodel(Sbwd,train_set).aic else Sbwd
print(Sboth)

{'residual_sugar:alcohol', 'fixed_acidity:sulphates', 'density:pH'}
```

### OLS Model with Forward Stepwise Terms

```
In [34]: > #fitting model based on optimal AIC stepwise regression
linear_formula = 'quality ~' + '+'.join(list(Sboth))
lm4 = smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 After AIC Stepwise: ', round(lm4.rsquared,3))
```

In Sample R2 After AIC Stepwise: 0.0

```
In [35]: > features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation with AIC Stepwise Regression ', CV_R2(features,line
print('Out of Sample MSE via Cross Validation with AIC Stepwise Regression ', CV_R2(features,line

Out of Sample R^2 via Cross Validation with AIC Stepwise Regression -0.0009112009929356901
Out of Sample MSE via Cross Validation with AIC Stepwise Regression 0.752974226749382
```

Since the forward regression only produced an optimal AIC with {'residual\_sugar:alcohol', 'fixed\_acidity:sulphates', 'density:pH'} terms, the MSE is decently high and the R-squared is near 0 because the coefficients are nearly 0. Indicating that this model is mostly a horizontal line.

### OLS Model with Backwards Stepwise Terms

```
In [37]: #fitting model based on optimal AIC stepwise regression via Stepwise Backwards
linear_formula = 'quality ~' + '+'.join(list(Sbwd))
lm5= smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 After AIC Stepwise Backward Regression: ', round(lm5.rsquared,3))
```

In Sample R2 After AIC Stepwise Backward Regression: 0.321

```
In [38]: features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation with AIC Stepwise Regression Backwards ', CV_R2(fea
print('Out of Sample MSE via Cross Validation with AIC Stepwise Regression Backwards ', CV_R2(fea
```

Out of Sample R^2 via Cross Validation with AIC Stepwise Regression Backwards 0.15321236823614456  
 Out of Sample MSE via Cross Validation with AIC Stepwise Regression Backwards 0.6344754312783343

Since the backward stepwise regression produced an optimal AIC with many terms, the MSE and r-squared out of sample did not improve. Indicating that this model is adding more noise than true correlations and overfitting the training set. The backwards stepwise regression contains these terms:

```
{'fixed_acidity:residual_sugar', 'total_sulfur_dioxide:pH', 'alcohol', 'residual_sugar:volatile_acidity', 'pH',
'volatile_acidity:chlorides', 'citric_acid:Color', 'residual_sugar:pH', 'fixed_acidity', 'residual_sugar:Color', 'Color',
'total_sulfur_dioxide:residual_sugar', 'chlorides', 'fixed_acidity:sulphates', 'density:Color', 'fixed_acidity:free_sulfur_dioxide',
'free_sulfur_dioxide:alcohol', 'free_sulfur_dioxide:volatile_acidity', 'free_sulfur_dioxide', 'density:alcohol',
'density:volatile_acidity', 'citric_acid:volatile_acidity', 'total_sulfur_dioxide:chlorides', 'sulphates', 'residual_sugar:alcohol',
'sulphates:alcohol', 'sulphates:chlorides', 'sulphates:volatile_acidity', 'fixed_acidity:alcohol', 'fixed_acidity:total_sulfur_dioxide',
'alcohol:Color', 'citric_acid:total_sulfur_dioxide', 'sulphates:pH', 'citric_acid:residual_sugar', 'total_sulfur_dioxide:density',
'fixed_acidity:density', 'citric_acid', 'fixed_acidity:volatile_acidity', 'total_sulfur_dioxide', 'density:pH',
'citric_acid:free_sulfur_dioxide', 'citric_acid:density', 'citric_acid:sulphates', 'sulphates:residual_sugar', 'residual_sugar',
'free_sulfur_dioxide:chlorides', 'volatile_acidity', 'citric_acid:chlorides', 'alcohol:pH', 'Color:chlorides',
'total_sulfur_dioxide:Color', 'density:chlorides', 'total_sulfur_dioxide:alcohol', 'residual_sugar:chlorides',
'free_sulfur_dioxide:density', 'total_sulfur_dioxide:volatile_acidity', 'fixed_acidity:chlorides', 'free_sulfur_dioxide:Color',
'free_sulfur_dioxide:sulphates', 'free_sulfur_dioxide:pH', 'sulphates:Color', 'free_sulfur_dioxide:residual_sugar',
'density:sulphates', 'alcohol:chlorides', 'pH:volatile_acidity'}
```

## Part H - Fit a Lasso model of 2-Way Interaction terms Report Estimates of Prediction Error



```
In [29]: from sklearn import linear_model, model_selection, tree, ensemble
from sklearn.linear_model import LassoCV

features = set(singlefeatures).union(set(interactions))
X = train_set.loc[:,list(set(df.columns)-{'quality'})]
#adding the interaction columns because Sklearn needs every column unlike stats models
for terms in interactions:
    first,second=terms.split(':')
    #creating a multiplication for the two columns as a two way interaction
    X[terms] = X[first]*X[second]

y = train_set.quality

lasso = LassoCV(cv=5, random_state=0).fit(X, y)
lassofeatures = list(zip(lasso.feature_names_in_,np.abs(lasso.coef_)>1e-10)) #getting the lasso f
lassofeatures = [x[0] for x in lassofeatures if x[1] == True] #filtering by the ones that are tru
lamb = lasso.alpha_
print('Lambda via CV Lasso ', lamb)
print('Lasso Features ', lassofeatures)
```

```
Lambda via CV Lasso  0.03914400323000001
Lasso Features  ['total_sulfur_dioxide', 'free_sulfur_dioxide:alcohol', 'free_sulfur_dioxide:p
H', 'free_sulfur_dioxide:Color', 'free_sulfur_dioxide:total_sulfur_dioxide', 'free_sulfur_dioxid
e:residual_sugar', 'free_sulfur_dioxide:fixed_acidity', 'alcohol:volatile_acidity', 'alcohol:sul
phates', 'alcohol:pH', 'alcohol:total_sulfur_dioxide', 'alcohol:residual_sugar', 'alcohol:fixed_
acidity', 'chlorides:total_sulfur_dioxide', 'volatile_acidity:total_sulfur_dioxide', 'sulphates:
total_sulfur_dioxide', 'pH:total_sulfur_dioxide', 'Color:total_sulfur_dioxide', 'total_sulfur_di
oxide:residual_sugar', 'total_sulfur_dioxide:fixed_acidity', 'residual_sugar:fixed_acidity']
```

```
In [30]: linear_formula = 'quality ~' + '+'.join(list(lassofeatures))
lm6= smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 After Lasso Feature Selection ', round(lm6.rsquared,3))
```

```
In Sample R2 After Lasso Feature Selection  0.319
```

```
In [31]: features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation with Lasso Features', CV_R2(features,linear_formula
print('Out of Sample MSE via Cross Validation with Lasso Features', CV_R2(features,linear_formula
```

```
Out of Sample R^2 via Cross Validation with Lasso Features 0.3009848094030164
Out of Sample MSE via Cross Validation with Lasso Features 0.5267910390150287
```

After performing Lasso for two-way interactions and single dimension features with 5-fold cv, the lambda is 0.039. The insample r-squared was 0.319 and the out of sample r\_squared is 0.3 with an MSE of 0.5216. This is the best performance we have seen so far. The Lasso terms includes ['total\_sulfur\_dioxide', 'free\_sulfur\_dioxide:alcohol', 'free\_sulfur\_dioxide:pH', 'free\_sulfur\_dioxide:Color', 'free\_sulfur\_dioxide:total\_sulfur\_dioxide', 'free\_sulfur\_dioxide:residual\_sugar', 'free\_sulfur\_dioxide:fixed\_acidity', 'alcohol:volatile\_acidity', 'alcohol:sulphates', 'alcohol:pH', 'alcohol:total\_sulfur\_dioxide', 'alcohol:residual\_sugar', 'alcohol:fixed\_acidity', 'chlorides:total\_sulfur\_dioxide', 'volatile\_acidity:total\_sulfur\_dioxide', 'sulphates:total\_sulfur\_dioxide', 'pH:total\_sulfur\_dioxide', 'Color:total\_sulfur\_dioxide', 'total\_sulfur\_dioxide:residual\_sugar', 'total\_sulfur\_dioxide:fixed\_acidity', 'residual\_sugar:fixed\_acidity']

## Part I - Pick favorite model and compare to the previous models

```
In [32]: linear_formula = 'quality ~' + '+'.join(list(lassofeatures) + singlefeatures)
lm7= smf.ols(linear_formula, data=train_set).fit()
print('In Sample R2 After Lasso + Single Feature Selection ', round(lm7.rsquared,3))
features = list(set(df.columns)-{'quality'})
print('Out of Sample R^2 via Cross Validation with Lasso Features', CV_R2(features,linear_formula
print('Out of Sample MSE via Cross Validation with Lasso Features', CV_R2(features,linear_formula
```

```
In Sample R2 After Lasso + Single Feature Selection  0.333
Out of Sample R^2 via Cross Validation with Lasso Features 0.3070809199544662
Out of Sample MSE via Cross Validation with Lasso Features 0.5221288590003825
```

This model that combines the Lasso 2-way interaction terms with all single features yields the highest out of sample r-squared

at 0.307 compare to previous models. The MSE for out of sample test errors in 0.522 which is marginally better than the LassoCV model.

## Part J - Create a 95% Wald predictive confidence for wine quality

```
In [48]: ▶ #Function that generates the Wald 95% CI and tests to see if true wine quality is in interval for
#Uses the model trained in part F for all single and two-way interaction terms
def generateWaldConfidenceIntervals(train_set, test_set):
    #Model from part F single and two way interactions
    linear_formula = 'quality ~' + '+' + '.join(list(set(df.columns)-{'quality'}))+['%s:%s'%v for v in
    lm3 = smf.ols(linear_formula, data=train_set).fit()
    #Gets the list of confidence intervals for the test_set records using Stat's Modles conf_int
    confidence_intervals = lm3.get_prediction(test_set).conf_int(obs=True, alpha=0.05)
    #Obtain the true wine quality
    true_quality = list(test_set.quality)
    correct = 0
    #Loop through the two lists and count whether the true wine quality is in the interval
    for i, interval in enumerate(confidence_intervals):
        if interval[0] <= true_quality[i] <= interval[1]:
            correct+=1
    #Returns the percentage of true wine quality
    return correct/len(true_quality)
correct = generateWaldConfidenceIntervals(train_set, test_set)
print('One run of WI and percentage of times contains true wine quality ', correct)
```

One run of WI and percentage of times contains true wine quality 0.9465597862391449

## Repeating Wald CI Experiment 100x over randomly generated train-test splits

```
In [47]: ▶ #running experiment 100 times
from random import sample
waldIntervalsCorrect = np.zeros(100)
for i in range(100):
    #Obtains random samples
    train_indices = list(sample(list(df.index), 5000))
    test_indices = list(set(list(df.index)) - set(train_indices))
    train_set = df.iloc[train_indices,: ]
    test_set = df.iloc[test_indices,: ]
    #Store percentage correct for each experiment run
    waldIntervalsCorrect[i] = generateWaldConfidenceIntervals(train_set, test_set)
print('Percentage WI contrains true quality', waldIntervalsCorrect.mean(), 'Standard Deviation', w
```

Percentage WI contrains true quality 0.9458984635938544 Standard Deviation 0.005453533066307316

## Part K - Repeat Part J but with Stepwise Regression

I am using forward stepwise regression because it is computationally faster than backwards and has the lowest AIC. After getting the Sfdw features I run an OLS model to get predictions on test set and determine if the actual wine quality lies in the interval.

```
In [55]: #Function that does the same steps as above but calls Forward Stepwise regression because in Part
#forward stepwise regression was chosen over backwards as it is computationally faster
def generateWaldConfidenceIntervalsStepWise(train_set, test_set, Sfdw):
    linear_formula = 'quality ~' + '+' .join(list(Sfdw))
    lm4 = smf.ols(linear_formula, data=train_set).fit()
    confidence_intervals = lm4.get_prediction(test_set).conf_int(obs=True, alpha=0.05)
    true_quality = list(test_set.quality)
    correct = 0
    for i, interval in enumerate(confidence_intervals):
        if interval[0] <= true_quality[i] <= interval[1]:
            correct+=1
    return correct/len(true_quality)

def fitmodel(S,data):
    return smf.ols('quality ~ '+'+' .join(S) if len(S)>0 else '1'), data=data).fit()

def forwardStepWise(singlefeatures, interactions, train_set):
    ## forward stepwise
    Sfdw = set()
    # features = set(singlefeatures)
    features = set(singlefeatures).union(set(interactions))
    while len(Sfdw)<len(features):
        f = max(features - Sfdw, key = lambda f: fitmodel(Sfdw.union({f}),train_set).aic)
        after = fitmodel(Sfdw.union({f}),train_set).aic
        before = fitmodel(Sfdw,train_set).aic
        if after > before:
            Sfdw = Sfdw.union({f})
        else:
            break
    return Sfdw

singlefeatures = list(set(df.columns)-{'quality'})
interactions = ['%s:%s'%v for v in itertools.combinations(singlefeatures,2)]
Sfdw = forwardStepWise(singlefeatures, interactions, train_set)
correct = generateWaldConfidenceIntervalsStepWise(train_set, test_set, Sfdw)
print('One run of Stepwise trained WI and percentage of times contains true wine quality ', corre

One run of Stepwise trained WI and percentage of times contains true wine quality 0.93319973279
89312
```

### Repeating Wald CI Experiment 10 over randomly generated train-test splits.

My computer was taking very long to run 100 experiments. I also use forward stepwise regression because the algorithm finds features much faster than backwards. After getting the Sfdw features I run an OLS model to get predictions on test set and determine if the actual wine quality lies in the interval.

```
In [56]: #running experiment 10 times
from random import sample
n = 10
waldIntervalsCorrect = np.zeros(n)
for i in range(n):
    train_indices = list(sample(list(df.index), 5000))
    test_indices = list(set(list(df.index)) - set(train_indices))
    train_set = df.iloc[train_indices,:]
    test_set = df.iloc[test_indices,:]
    Sfdw = forwardStepWise(singlefeatures, interactions, train_set)
    waldIntervalsCorrect[i] = generateWaldConfidenceIntervalsStepWise(train_set, test_set, Sfdw)
print('Percentage WI contrains true quality', waldIntervalsCorrect.mean(), 'Standard Deviation',w

Percentage WI contrains true quality 0.9345357381429527 Standard Deviation 0.007744222710054484
```

## Part L - Use Bootstrap to generate confidence intervals

Approach: Since the question was not entirely clear I assumed that the goal is to use bootstrap to get a pivotal CI for the mean of wine quality. I am using the stepwise regression from part K to fit a model and then calculate the mean of the predicted wine quality and testing if its in the bootstrap CI. I run this experiment 25 times of speed.

```
In [39]: from arch.bootstrap import IIDBootstrap

In [69]: def generateBootStrapCI(train_set, test_set, Sfdw):
    #Fit a model based on Sfdw features
    linear_formula = 'quality ~' + '+'.join(list(Sfdw))
    lm4 = smf.ols(linear_formula, data=train_set).fit()
    #Theta function set to the mean
    thetahn = lambda z: np.mean(z)
    #Get the predictions from the model
    predictions = list(lm4.predict(test_set))
    thetahn = thetahn(predictions)
    #Using arch bootstrap function to generate pivotal CI
    CI = IIDBootstrap(test_set.quality).conf_int(
        thetahn,
        1000,
        method='basic')
    #return 1 if the the is in the interval or not
    if CI[0] <= thetahn <= CI[1]:
        return 1
    else:
        return 0

from random import sample
n = 25
bootIntervalCorrect = np.zeros(n)
for i in range(n):
    #resample the train and test sets
    train_indices = list(sample(list(df.index), 5000))
    test_indices = list(set(list(df.index)) - set(train_indices))
    train_set = df.iloc[train_indices,:]
    test_set = df.iloc[test_indices,:]
    #get the forward step wise features
    Sfdw = forwardStepWise(singlefeatures, interactions, train_set)
    #Calculate if the mean of the predictions is in the pivotal bootstrap CI
    bootIntervalCorrect[i] = generateBootStrapCI(train_set, test_set, Sfdw)

print('Percentage Bootstrap CI for mean constrains true quality', bootIntervalCorrect.mean(), 'Sta
844
```

Percentage Bootstrap CI for mean constrains true quality 0.88 Standard Deviation 0.32496153618543

## Part M - Repeat Pivotal Bootstrap for Mean using Lasso Model

Repeated the general steps in part K except train a 5-Fold Lasso CV model. Counting the number of times the predicted wine quality mean lies in the Bootstrap CI generated.

```

In [70]: > from sklearn import linear_model, model_selection, tree, ensemble
> from sklearn.linear_model import LassoCV

def generateBootStrapCILasso(train_set, test_set, lassofeatures):
    #Fit a model based on Sfdw features
    linear_formula = 'quality ~' + '+'.join(list(lassofeatures))
    lm6= smf.ols(linear_formula, data=train_set).fit()
    #Theta function set to the mean
    thetahn = lambda z: np.mean(z)
    #Get the predictions from the model
    predictions = list(lm6.predict(test_set))
    thetahn = thetahn(predictions)
    #Using arch bootstrap function to generate pivotal CI
    CI = IIDBootstrap(test_set.quality).conf_int(
        thetahn,
        1000,
        method='basic')
    #return 1 if the mean of the predictions is in the bootstrap pivotal interval
    if CI[0] <= thetahn <= CI[1]:
        return 1
    else:
        return 0

def lassoFeatures(train_set,df,singlefeatures,interactions):
    features = set(singlefeatures).union(set(interactions))
    X = train_set.loc[:,list(set(df.columns)-{'quality'})]
    #adding the interaction columns because Sklearn needs every column unlike stats models
    for terms in interactions:
        first,second=terms.split(':')
        #creating a multiplication for the two columns as a two way interaction
        X[terms] = X[first]*X[second]

    y = train_set.quality
    lasso = LassoCV(cv=5, random_state=0).fit(X, y)
    lassofeatures = list(zip(lasso.feature_names_in_,np.abs(lasso.coef_)>1e-10)) #getting the las
    lassofeatures = [x[0] for x in lassofeatures if x[1] == True] #filtering by the ones that are
    return lassofeatures

#running Lasso Experiment
from random import sample
n = 25
bootIntervalCorrect = np.zeros(n)
for i in range(n):
    #resample the train and test sets
    train_indices = list(sample(list(df.index), 5000))
    test_indices = list(set(list(df.index)) - set(train_indices))
    train_set = df.iloc[train_indices,:]
    test_set = df.iloc[test_indices,:]
    #get the forward step wise features
    lasso = lassoFeatures(train_set,df,singlefeatures,interactions)
    #Calculate if the mean of the predictions is in the pivotal bootstrap CI
    bootIntervalCorrect[i] = generateBootStrapCILasso(train_set, test_set, lasso)

print('Percentage Lasso Bootstrap CI for mean constrains true quality', bootIntervalCorrect.mean())

```

Percentage Lasso Bootstrap CI for mean constrains true quality 0.96 Standard Deviation 0.19595917 942265428

## Problem 5

```

In [10]: > ### Import data
> import pandas as pd
> import statsmodels.api as sm
df2 = pd.DataFrame(sm.datasets.star98.load().data)

```

## Part A - Simple OLS Student Math Scores vs Per-Pupil Spending

Run a simple regression to investigate whether student outcomes, as the fraction of students above the national mean, are affected by per-pupil spending (hint: if A and B are two column names, you can use  $I(A/(A+B)) \sim \dots$  as a specification of the response variable in statsmodels formula). Interpret the results – do you think there is a causal effect? Why or why not?

```
In [18]: import statsmodels.formula.api as smf
response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features = 'PERSPENK'
linear_formula = response + features
lm = smf.ols(linear_formula, data=df2).fit()
perspenkCoeff = lm.params['PERSPENK']
print('PERSPENK COEFF: ', perspenkCoeff)
lm.summary()
```

PERSPENK COEFF: 0.04917284368069601

Out[18]: OLS Regression Results

Dep. Variable:	I((NABOVE) / (NABOVE + NBELOW))		R-squared:	0.025		
Model:	OLS		Adj. R-squared:	0.022		
Method:	Least Squares		F-statistic:	7.665		
Date:	Mon, 18 Apr 2022		Prob (F-statistic):	0.00598		
Time:	10:54:41		Log-Likelihood:	86.908		
No. Observations:	303		AIC:	-169.8		
Df Residuals:	301		BIC:	-162.4		
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.2245	0.077	2.899	0.004	0.072	0.377
PERSPENK	0.0492	0.018	2.769	0.006	0.014	0.084
Omnibus:	8.083	Durbin-Watson:	1.646			
Prob(Omnibus):	0.018	Jarque-Bera (JB):	6.923			
Skew:	0.293	Prob(JB):	0.0314			
Kurtosis:	2.546	Cond. No.	33.9			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

### Interpretation of Simple OLS PERSPENK

Looking at the OLS Regression Results we can see that the  $R^2$  is very small 0.025 which means that the variation of Y is not well explained by the covariate set 'PERSPENK'. While this coefficient is significant and positively correlated with math score we cannot conclude causal effects as we have not enforced randomization of assignment to treatment or used ignorability principles to compare like-to-like counties.

## Part B - Simple OLS Student Math Scores vs Pupil-Teacher Ratio

```
In [12]: response2 = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features2 = 'PTRATIO'
linear_formula2 = response2 + features2
lm2 = smf.ols(linear_formula2, data=df2).fit()
ptratioCoeff = lm2.params['PTRATIO']
print('PTRATIO COEFF: ', ptratioCoeff)
lm2.summary()
```

PTRATIO COEFF: -0.014078019457326896

Out[12]:

OLS Regression Results

<b>Dep. Variable:</b>	I((NABOVE) / (NABOVE + NBELOW))	<b>R-squared:</b>	0.029
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.025
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	8.853
<b>Date:</b>	Fri, 15 Apr 2022	<b>Prob (F-statistic):</b>	0.00316
<b>Time:</b>	19:03:14	<b>Log-Likelihood:</b>	87.490
<b>No. Observations:</b>	303	<b>AIC:</b>	-171.0
<b>Df Residuals:</b>	301	<b>BIC:</b>	-163.6
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	0.7532	0.107	7.052	0.000	0.543	0.963
<b>PTRATIO</b>	-0.0141	0.005	-2.975	0.003	-0.023	-0.005

<b>Omnibus:</b>	5.689	<b>Durbin-Watson:</b>	1.643
<b>Prob(Omnibus):</b>	0.058	<b>Jarque-Bera (JB):</b>	5.583
<b>Skew:</b>	0.294	<b>Prob(JB):</b>	0.0613
<b>Kurtosis:</b>	2.689	<b>Cond. No.</b>	231.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

### Interpretation of OLS PTRatio

Looking at the OLS Regression Results we can see that the  $R^2$  is very small 0.020 which means that the variation of Y is not well explained by the covariate set 'PTRATIO' alone. While this coefficient is significant and negatively correlated with math score we cannot conclude causal effects as we have not enforced randomization of assignment to treatment or used ignorability principles to compare like-to-like counties.

## Part C - Adding Features to the OLS Model

For each of part (a) and part (b), try including additional variables in the following order:

Average teacher experience

Average teacher experience and student race

Finally, run a regression against all variables (but exclude the interactions terms). For each case, report the change in the effect coefficient for each model.

```
In [13]: # Case 1 Part A with AVYRSEXP
response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features = 'PERSPENK' + '+' + 'AVYRSEXP'
linear_formula = response + features
lm = smf.ols(linear_formula, data=df2).fit()
perspenkCoeff1 = lm.params['PERSPENK']
print('PERSPENK COEFF CASE 1: ', perspenkCoeff1)
print('PERSPENK Coefficient Difference after adding AVYRSEXP ', round(perspenkCoeff1 - perspenkCo
```

```
PERSPENK COEFF CASE 1: 0.028124399361265873
PERSPENK Coefficient Difference after adding AVYRSEXP -0.021048
```

```
In [14]: # Case 2 Part A with AVYRSEXP + Student Race
response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features = ['PERSPENK', 'AVYRSEXP', 'PERASIAN', 'PERBLACK', 'PERHISP']
linear_formula = response + '+'.join(features)
lm = smf.ols(linear_formula, data=df2).fit()
perspenkCoeff2 = lm.params['PERSPENK']
print('PERSPENK COEFF CASE 2: ', perspenkCoeff2)
print('PERSPENK Coefficient Difference after adding AVYRSEXP + STUDENT RACE ', round(perspenkCoef
```

```
PERSPENK COEFF CASE 2: -0.002637582935566419
PERSPENK Coefficient Difference after adding AVYRSEXP + STUDENT RACE -0.05181
```

```
In [15]: # Case 1 Part B with AVYRSEXP
response2 = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features2 = 'PTRATIO' + '+' + 'AVYRSEXP'
linear_formula2 = response2 + features2
lm2 = smf.ols(linear_formula2, data=df2).fit()
ptratioCoeff1 = lm2.params['PTRATIO']
print('PTRATIO COEFF CASE 1: ', ptratioCoeff1)
print('PTRATIO Coefficient Difference after adding AVYRSEXP ', round(ptratioCoeff1 - ptratioCoef
```

```
PTRATIO COEFF CASE 1: -0.01480897162953032
PTRATIO Coefficient Difference after adding AVYRSEXP -0.000731
```

```
In [16]: # Case 2 Part B with AVYRSEXP + Student Race
response2 = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features2 = ['PTRATIO', 'AVYRSEXP', 'PERASIAN', 'PERBLACK', 'PERHISP']
linear_formula2 = response2 + '+'.join(features2)
lm2 = smf.ols(linear_formula2, data=df2).fit()
ptratioCoeff2 = lm2.params['PTRATIO']
print('PTRATIO COEFF CASE 2: ', ptratioCoeff2)
print('PTRATIO Coefficient Difference after adding AVYRSEXP + Student Race ', round(ptratioCoeff2
```

```
PTRATIO COEFF CASE 2: 0.002653121266625309
PTRATIO Coefficient Difference after adding AVYRSEXP + Student Race 0.016731
```

```
In [17]: # Case 3 - Finally, run a regression against all variables (but exclude the interactions terms).
features = list(df2.columns)
single_features = [x for x in features if '_' not in x][2:]
response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
linear_formula = response + '+'.join(single_features)
lm = smf.ols(linear_formula, data=df2).fit()
perspenkCoeff3 = lm.params['PERSPENK']
print('PERSPENK COEFF CASE 3: ', perspenkCoeff3)
ptratioCoeff3 = lm.params['PTRATIO']
print('PTRATIO COEFF CASE 3: ', ptratioCoeff3)
print('PERSPENK Coefficient Difference after All Single Features ', round(perspenkCoeff3 - perspe
print('PTRATIO Coefficient Difference after All Single Features', round(ptratioCoeff3 - ptratioCo
```

```
PERSPENK COEFF CASE 3: 0.01024358094521226
PTRATIO COEFF CASE 3: -0.000980990866091552
PERSPENK Coefficient Difference after All Single Features -0.038929
PTRATIO Coefficient Difference after All Single Features 0.013097
```



### ***Conclusion of Coefficient Analysis***

As we added more terms for Model A the PERSPENK coefficient had a smaller effect than the original model. For Model B the PTRATIO started from a negative effect and grew closer to 0 as we added terms. For both experiments the addition of Student Race covariates made the coefficients flip (PERSPENK positive to negative and PTRATIO negative to positive).

### **Part D - Experimenting with Model and Interaction Terms and Observing changes in coefficients**

Experiment with the data to come up with your favorite model for student outcomes. Try adding the interaction terms. What do you observe? How do the coefficients change?

### **Running a model against full feature set**

```
In [51]: response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features = df2.columns
#Removing features that are not significant
features = set(features) #- {'NABOVE', 'NBELOW', 'PERMINTE', 'AVYRSEXP', 'AVSALK', 'PCTAF', 'PCTCHRT', '
interactions = '+'.join(list(interaction))
linear_formula = response + '+'.join(list(features))
lm9 = smf.ols(linear_formula, data=df2).fit()
lm9.summary()
```

Out[51]: OLS Regression Results

Dep. Variable:	I((NABOVE) / (NABOVE + NBELOW))			R-squared:		0.830	
Model:	OLS			Adj. R-squared:		0.817	
Method:	Least Squares			F-statistic:		62.15	
Date:	Mon, 18 Apr 2022			Prob (F-statistic):		3.28e-94	
Time:	11:15:36			Log-Likelihood:		351.58	
No. Observations:	303			AIC:		-657.2	
Df Residuals:	280			BIC:		-571.7	
Df Model:	22						
Covariance Type:	nonrobust						
		coef	std err	t	P> t	[0.025	0.975]
Intercept	0.5039	0.916	0.550	0.583	-1.299	2.306	
PERASIAN	0.0021	0.001	3.390	0.001	0.001	0.003	
PERSPENK	0.0036	0.155	0.023	0.982	-0.302	0.309	
NABOVE	9.986e-05	1.95e-05	5.124	0.000	6.15e-05	0.000	
PERMINTE_AVYRSEXP	-0.0012	0.002	-0.480	0.632	-0.006	0.004	
PERBLACK	-0.0042	0.001	-5.760	0.000	-0.006	-0.003	
AVSALK	-0.0019	0.011	-0.162	0.872	-0.024	0.021	
PCTCHRT	-0.0008	0.001	-0.986	0.325	-0.002	0.001	
PERSPEN_PCTAF	0.0004	0.004	0.108	0.914	-0.007	0.007	
LOWINC	-0.0042	0.000	-10.934	0.000	-0.005	-0.003	
PCTAF	0.0045	0.017	0.262	0.793	-0.029	0.038	
AVYRSEXP_AVSAL	0.0004	0.001	0.453	0.651	-0.001	0.002	
PTRATIO	0.0170	0.034	0.506	0.613	-0.049	0.083	
PERMINTE_AVSAL	-0.0004	0.001	-0.743	0.458	-0.002	0.001	
PERMINTE	0.0267	0.035	0.765	0.445	-0.042	0.095	
PERSPEN_PTRATIO_PCTAF	5.647e-05	0.000	0.321	0.749	-0.000	0.000	
PCTYRRND	-0.0004	0.000	-1.819	0.070	-0.001	3.42e-05	
NBELOW	-3.456e-05	8.65e-06	-3.995	0.000	-5.16e-05	-1.75e-05	
PERHISP	-0.0019	0.000	-5.280	0.000	-0.003	-0.001	
PERMINTE_AVYRSEXP_AVSAL	1.98e-05	4.1e-05	0.483	0.629	-6.09e-05	0.000	
PTRATIO_PCTAF	-0.0005	0.001	-0.583	0.560	-0.002	0.001	
PERSPEN_PTRATIO	-0.0023	0.008	-0.307	0.759	-0.017	0.013	
AVYRSEXP	-0.0168	0.046	-0.367	0.714	-0.107	0.073	
Omnibus:	5.088	Durbin-Watson:	1.927				
Prob(Omnibus):	0.079	Jarque-Bera (JB):	7.017				

<b>Skew:</b>	-0.019	<b>Prob(JB):</b>	0.0299
<b>Kurtosis:</b>	3.745	<b>Cond. No.</b>	3.56e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.56e+06. This might indicate that there are strong multicollinearity or other numerical problems.

**Running a model after removing insignificant features based on p-val**

```
In [52]: response = 'I((NABOVE)/(NABOVE+NBELOW)) ~ '
features = df2.columns
#Removing features that are not significant
features = set(features) - {'NABOVE', 'NBELOW', 'PERMINTE', 'AVYRSEXP', 'AVSALK', 'PCTAF', 'PCTCHRT', 'P
linear_formula = response + '+'.join(list(features))
lm10 = smf.ols(linear_formula, data=df2).fit()
lm10.summary()
```

Out[52]: OLS Regression Results

Dep. Variable:	I((NABOVE) / (NABOVE + NBELOW))		R-squared:		0.808		
Model:	OLS		Adj. R-squared:		0.800		
Method:	Least Squares		F-statistic:		101.5		
Date:	Mon, 18 Apr 2022		Prob (F-statistic):		3.07e-96		
Time:	11:28:57		Log-Likelihood:		332.88		
No. Observations:	303		AIC:		-639.8		
Df Residuals:	290		BIC:		-591.5		
Df Model:	12						
Covariance Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]	
Intercept	0.5240	0.105	4.968	0.000	0.316	0.732	
PERMINTE_AVYRSEXP	0.0005	0.000	1.756	0.080	-5.88e-05	0.001	
PERBLACK	-0.0045	0.001	-6.428	0.000	-0.006	-0.003	
PCTYRRND	-0.0004	0.000	-1.914	0.057	-0.001	1.24e-05	
PERASIAN	0.0032	0.001	5.316	0.000	0.002	0.004	
PERSPEN_PCTAF	0.0013	0.000	4.191	0.000	0.001	0.002	
PTRATIO	0.0093	0.004	2.308	0.022	0.001	0.017	
LOWINC	-0.0046	0.000	-11.840	0.000	-0.005	-0.004	
PERHISP	-0.0021	0.000	-5.729	0.000	-0.003	-0.001	
PERSPENK	-0.0349	0.017	-2.048	0.042	-0.068	-0.001	
PERMINTE_AVYRSEXP_AVSAL	-7.327e-06	4.35e-06	-1.683	0.093	-1.59e-05	1.24e-06	
PTRATIO_PCTAF	-0.0002	6.13e-05	-3.313	0.001	-0.000	-8.25e-05	
AVYRSEXP_AVSAL	9.579e-05	4.65e-05	2.060	0.040	4.25e-06	0.000	
Omnibus:	8.042	Durbin-Watson:	1.879				
Prob(Omnibus):	0.018	Jarque-Bera (JB):	11.211				
Skew:	-0.187	Prob(JB):	0.00368				
Kurtosis:	3.865	Cond. No.	3.86e+05				

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.86e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Based on the data, I chose a model where p-value indicated the term was significant whereas 'NABOVE', 'NBELOW', 'PERMINTE', 'AVYRSEXP', 'AVSALK', 'PCTAF', 'PCTCHRT', 'PERSPEN\_PTRATIO', 'PERSPEN\_PTRATIO\_PCTAF' terms were not.

Model Coefficients after experimenting with data:

- Intercept 0.524040
- PERMINTE\_AVYRSEXP 0.000486
- PERBLACK -0.004522
- PCTYRRND -0.000441
- PERASIAN 0.003245
- PERSPEN\_PCTAF 0.001301
- PTRATIO 0.009275
- LOWINC -0.004595
- PERHISP -0.002104
- PERSPENK -0.034890
- PERMINTE\_AVYRSEXP\_AVSAL -0.000007
- PTRATIO\_PCTAF -0.000203
- AVYRSEXP\_AVSAL 0.000096

The resulting model an in sample r-squared of 0.808 and the PERSPENK = -0.034 and PTRATIO = 0.009. With all features PERSPENK = 0.0036 and PTRATIO = 0.0170. Removing coefficients made PTRATIO smaller and PERSPENK negative and in absolute terms bigger.



## Part E - Interpreting Causal Effects on student outcomes

Since PTRATIO has a positive confidence interval and PERSPENK has a negative confidence interval we have 95% confidence that the true coefficient for pupil-teacher ratio is slightly positive. Similarly we are 95% confident the per-pupil spending has a negative effect on student outcomes. We can't say this is causal because there is not enforced randomization of treatment and measured outcomes like in an A/B test. Further investigation can be done because in this analysis these coefficients are near to 0 and switch signs depending on the number of features used in the model.